

Varun Sadashive Gowda

[linkedin.com/in/varunsg11/](https://www.linkedin.com/in/varunsg11/) | github.com/varunsg11 | varunsg.dev | varunsg118@gmail.com | College Station, Texas

Professional Summary

Incoming MCS student at Texas A&M with 3 years of industry experience building transformer-based LLM agents, RAG pipelines, and MLOps workflows. Architected multi-agent systems using CrewAI, LangChain, and LangGraph that reduced support effort by 30%, with hands-on experience in fine-tuning (PEFT/LoRA), model evaluation, vector search, and CI/CD for AI systems.

Education

Master's in Computer Science, Texas A&M University

Aug 2026 to May 2028

Course work - 629 Analysis of Algorithms, 636 Deep Learning, 670 Information Storage and Retrieval

Bachelor of Engineering, Electronics and Communication, JSSSTU

Jul 2019 to Jul 2023

Course work - Computer Networks, Operating system, Cryptography

8.94/10 CGPA (3.58/4.0)

Skills & Competencies

Languages - Python, ABAP, Java, C, C++, SQL, YAML, Go

Technical skills - Multi-Agent System Design & Orchestration, LLM Integration & Tooling (APIs, Function Calling, Automation), Transformer Architecture & Fine-tuning (PEFT/LoRA), Prompt Engineering, Generative AI, Machine Learning, Natural Language Processing (NLP), MLOps, Vector Databases & Semantic Search, Statistical Modeling

Tools/Platform - PyTorch, TensorFlow, Hugging Face, LangChain, LangGraph, CrewAI, scikit-learn, Pandas, NumPy, Jupyter, MLflow, Databricks, OpenAI API, vLLM, Docker, Kubernetes, GitHub, HANA DB

Experience

SAP Labs, Bengaluru: Agentic AI Developer

Sep 2025 to Aug 2026

- Designed and deployed multi-agent workflows using CrewAI, LangChain, and LangGraph reducing support engineers effort by 30% through automated task coordination and LLM-driven decision-making.
- Developed custom agent tools for data extraction, solution analysis, document parsing (PDF/CSV), and automated reporting, including an OCR-based image extraction pipeline for structured knowledge documents.
- Developed and maintained MCP (Model Context Protocol) server to enable seamless integration between the data warehouse and AI agent repositories, standardizing tool-calling interfaces across services.
- Built a RAG pipeline on Databricks using transformer-based embeddings for semantic chunking and vector indexing in SAP HANA DB, enabling sub-second retrieval across 50K+ enterprise documents for LLM context grounding.
- Built MLOps CI/CD pipelines (GitHub Actions) for the projects repository, including automated Databricks job orchestration for data processing workflows.
- Led Docker containerization efforts, cutting image size by 67% (6 GB to 2 GB) to streamline builds and deployments.

SAP Labs, Bengaluru: Associate Solution Support Engineer

Feb 2023 to Sep 2025

- Analyzed complex problems and resolved 600+ critical issues for SAP customers through incidents.
- Mentored junior engineers and external hires, enhancing team productivity by 20% through code/issue reviews and technical workshops on advanced topics of Transportation Management.
- Built and deployed two internal JS extensions used by 1,000+ support engineers worldwide.

Projects

Text Independent Speaker Recognition and Classification using KNN Algorithm ([Publication](#))

Conducted research on speaker recognition techniques, achieving a 96.97% accuracy using a double distance measurement method, surpassing the 84.85% accuracy of K-Nearest Neighbors (KNN).

Tech: Python, scikit-learn, NumPy, Audio Feature Extraction (MFCC), Model Evaluation

Real Time Theft Detection Using YOLOv5 Object Detection Model ([Publication](#))

Developed an automated theft detection system using the YOLOv5 object detection model, enabling real-time identification of objects and theft events with alerts to owners and authorities.

Tech: Python, PyTorch, YOLOv5, OpenCV, Real-time Inference, Model Benchmarking

Certifications: [SAP Generative AI Developer](#) | [Agentic AI Fundamentals](#) | [Build AI Agents & Chatbots with LangGraph](#)